

DS505: INTRODUCTION TO DEEP LEARNING

P04: Convolutional Neural Network (CNN)

```
In [1]: from PIL import Image, ImageFilter
```


```
In [3]: image = Image.open('dobermann.jpg').convert('RGB')
```

```
In [4]: display(image)
```



```
In [5]: filtered = image.filter(ImageFilter.Kernel(size=(3,3),  
                                                kernel=[-1,-1,-1,-1,8,-1,-1,-1,-1],  
                                                scale=1))
```

```
In [6]: display(filtered)
```



```
In [7]: import tensorflow as tf
        from tensorflow.keras.utils import to_categorical
        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
```

```
In [8]: mnist = tf.keras.datasets.mnist
```

```
In [9]: (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz> (<https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>)
11490434/11490434 [=====] - 4s 0us/step

```
In [10]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[10]: ((60000, 28, 28), (60000,), (10000, 28, 28), (10000,))
```

```
In [11]: index = 2
img = Image.fromarray(x_train[index])
print(y_train[index])
display(img)
```

4



```
In [12]: x_train = x_train / 255.0
x_test = x_test / 255.0
```

```
In [13]: y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

```
In [14]: y_train.shape
```

Out[14]: (60000, 10)

```
In [15]: x_train.shape
```

Out[15]: (60000, 28, 28)

```
In [16]: x_train = x_train.reshape(x_train.shape[0], x_train.shape[1], x_train.shape[2], 1)
x_test = x_test.reshape(x_test.shape[0], x_test.shape[1], x_test.shape[2], 1)
```

```
In [17]: x_train.shape
```

Out[17]: (60000, 28, 28, 1)

```
In [18]: model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(tf.keras.layers.MaxPooling2D(pool_size=(2, 2), input_shape=(4, 4, 1)))
#model.addMaxPooling2D((pool_size=(2, 2), strides=(2, 2), padding='valid'))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(.5))
model.add(Dense(10, activation='softmax'))
```

```
In [19]: print(model.summary())
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
flatten (Flatten)	(None, 5408)	0
dense (Dense)	(None, 128)	692352
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1290

=====

Total params: 693,962
Trainable params: 693,962
Non-trainable params: 0

None

```
In [20]: model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10)

model.evaluate(x_test, y_test)
```

```
Epoch 1/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.2710 - accuracy: 0.9179
Epoch 2/10
1875/1875 [=====] - 32s 17ms/step - loss: 0.1081 - accuracy: 0.9676
Epoch 3/10
1875/1875 [=====] - 30s 16ms/step - loss: 0.0828 - accuracy: 0.9754
Epoch 4/10
1875/1875 [=====] - 28s 15ms/step - loss: 0.0644 - accuracy: 0.9806
Epoch 5/10
1875/1875 [=====] - 24s 13ms/step - loss: 0.0551 - accuracy: 0.9830
Epoch 6/10
1875/1875 [=====] - 27s 14ms/step - loss: 0.0461 - accuracy: 0.9850
Epoch 7/10
1875/1875 [=====] - 30s 16ms/step - loss: 0.0414 - accuracy: 0.9866
Epoch 8/10
1875/1875 [=====] - 31s 16ms/step - loss: 0.0364 - accuracy: 0.9883
Epoch 9/10
1875/1875 [=====] - 31s 16ms/step - loss: 0.0325 - accuracy: 0.9898
Epoch 10/10
1875/1875 [=====] - 30s 16ms/step - loss: 0.0304 - accuracy: 0.9902
313/313 [=====] - 2s 5ms/step - loss: 0.0427 - accuracy: 0.9878
```

```
Out[20]: [0.042684439569711685, 0.9878000020980835]
```

```
In [ ]:
```